

29-11-2025

Project RAG | AI Coach for AWS Training

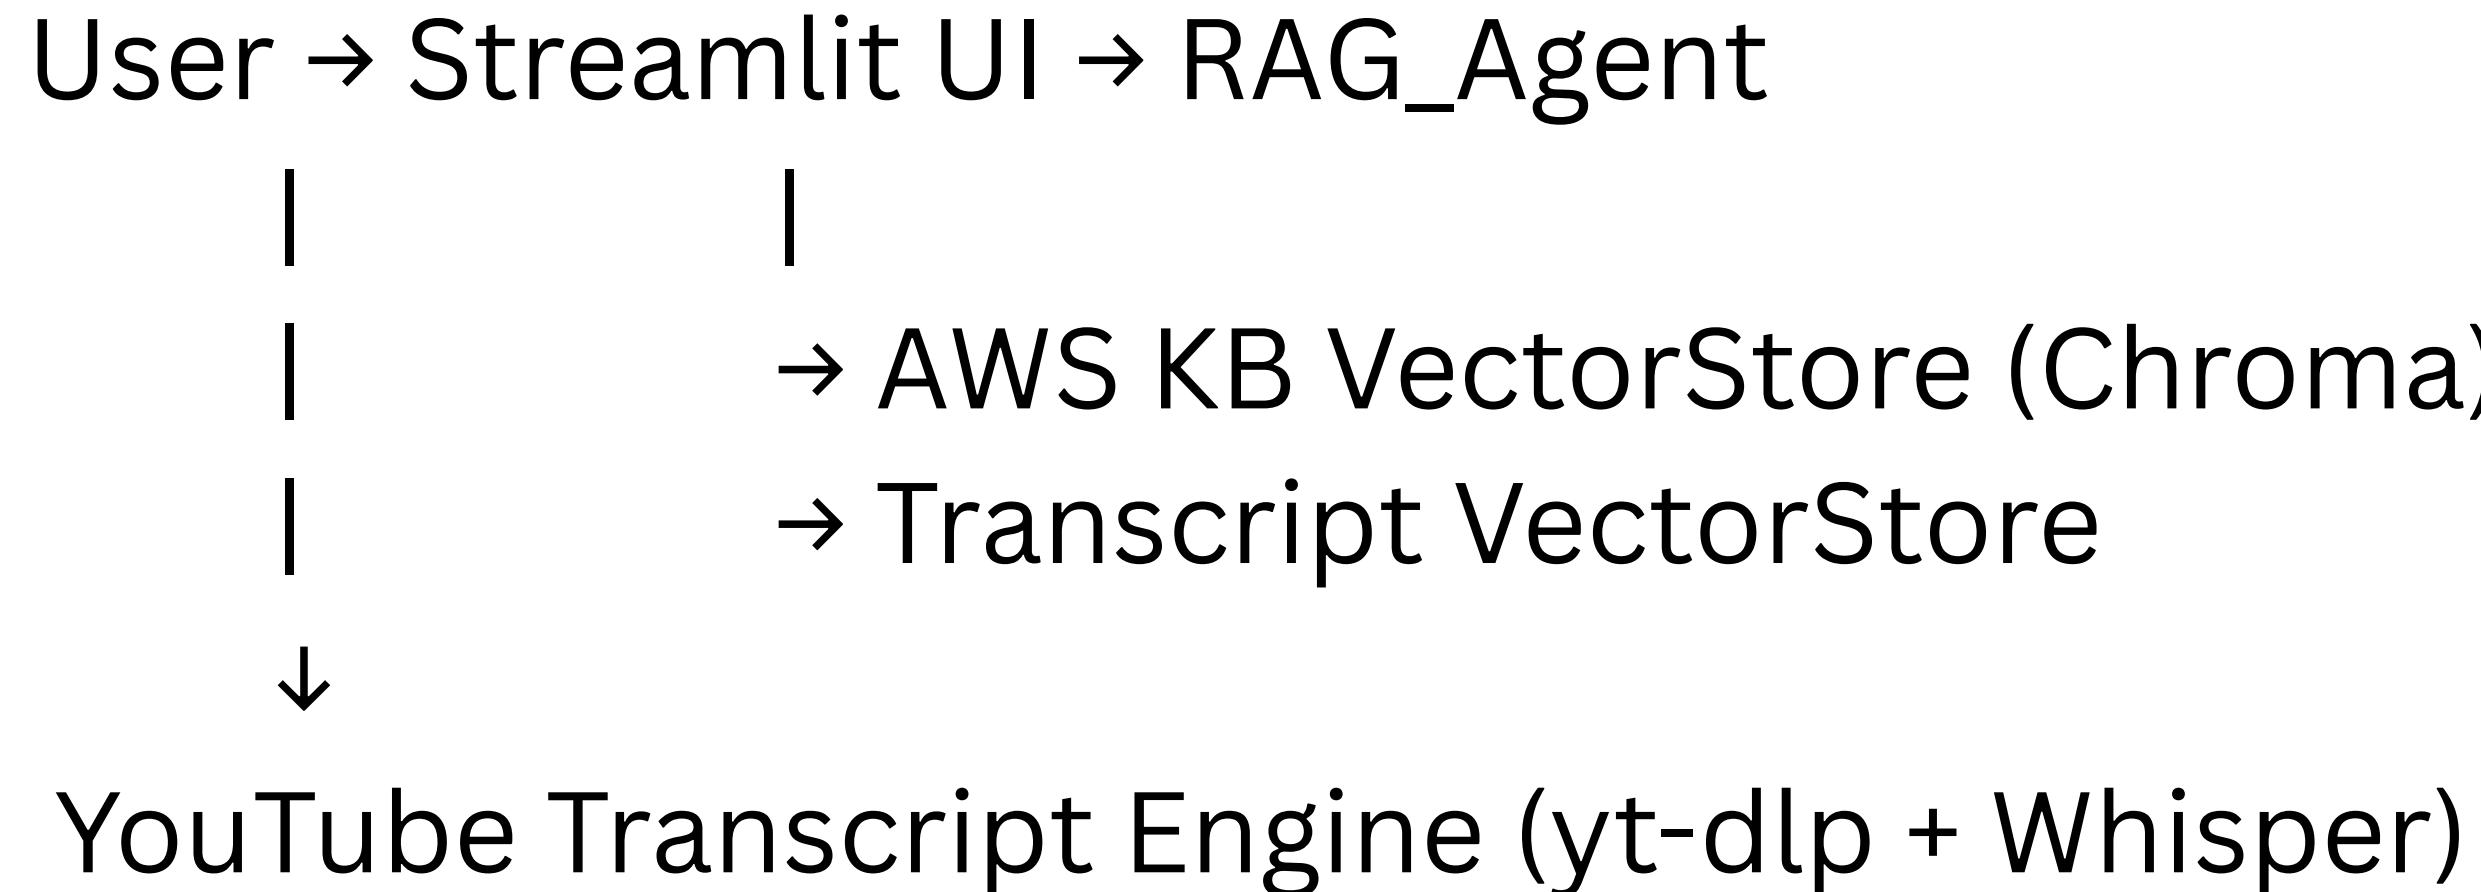
What Problem Are We Solving?

- AWS learning can be overwhelming for beginners
- Hard to find structured explanations, quizzes, or code examples
- No unified tool that can teach, quiz, generate code, and learn new topics from YouTube
- Why not use other AI tools??

What This System Can Do

- Explain any AWS topic
- Run quizzes (MCQs) with scoring + feedback
- Generate AWS code snippets (Python/Boto3)
- Learn new topics automatically from YouTube videos
- Store the knowledge in its own knowledge base (self-learning)

High-Level System Architecture



Key Components

- main.py – Controls logic & routing
- RAG_Agent.py – Core intelligence (teach, quiz, code)
- aws_info.py – Handles KB updates + vectorstore versioning
- NLP.py – Cleans and chunks transcripts
- transcript_extraction.py – Pulls YouTube subtitles / Whisper
- vector_store.py – Handles transcript embeddings
- app.py – The Streamlit user interface

Knowledge Base

- aws_info.py manages all AWS Knowledge Base operations:
- Loads + cleans raw KB text
- Splits into semantic chunks using spaCy
- Builds vectorstore versions (no overwriting)
- Stores meta.json with KB hash → auto-build new version on changes
- Provides retriever to RAG_Agent
- Computes relevance score to see if topic is already known

How RAG Checks Similarity & Answers Questions

- User asks: “What is AWS Glue?”
- RAG pipeline loads latest vectorstore version
- Computes similarity between query embedding and KB chunks
- If relevance score \leq threshold \rightarrow topic found \rightarrow **have to fix**
- Retrieves top chunks from `aws_knowledge_base.txt`
- LLM uses ONLY retrieved context to answer
- Prevents hallucinations by restricting to AWS KB content

User Query → Embedding → Chroma Similarity → Top Chunks → LLM Answer

RAG - Brain of the system

- Query normalization (adds AWS prefix, strips quiz words)
- Intent detection (teach / quiz / code)
- Retrieval selection (AWS KB or transcript vectorstore)
- TEACH mode → context-based explanation
- QUIZ mode → LLM-generated MCQs + evaluation
- CODE mode → generate Boto3/Python code with explanation

User Query → Intent → Normalize → Retrieve → LLM Prompt → Output

How the Bot Automatically Chooses Teach / Quiz / Code

Intent Classification Logic:

- If query contains quiz words → QUIZ mode
- If query contains code-related terms → CODE mode
- Otherwise → TEACH mode

Example Detection:

- “quiz on S3” → QUIZ
- “give python code for S3 upload” → CODE
- “explain AWS Glue” → TEACH

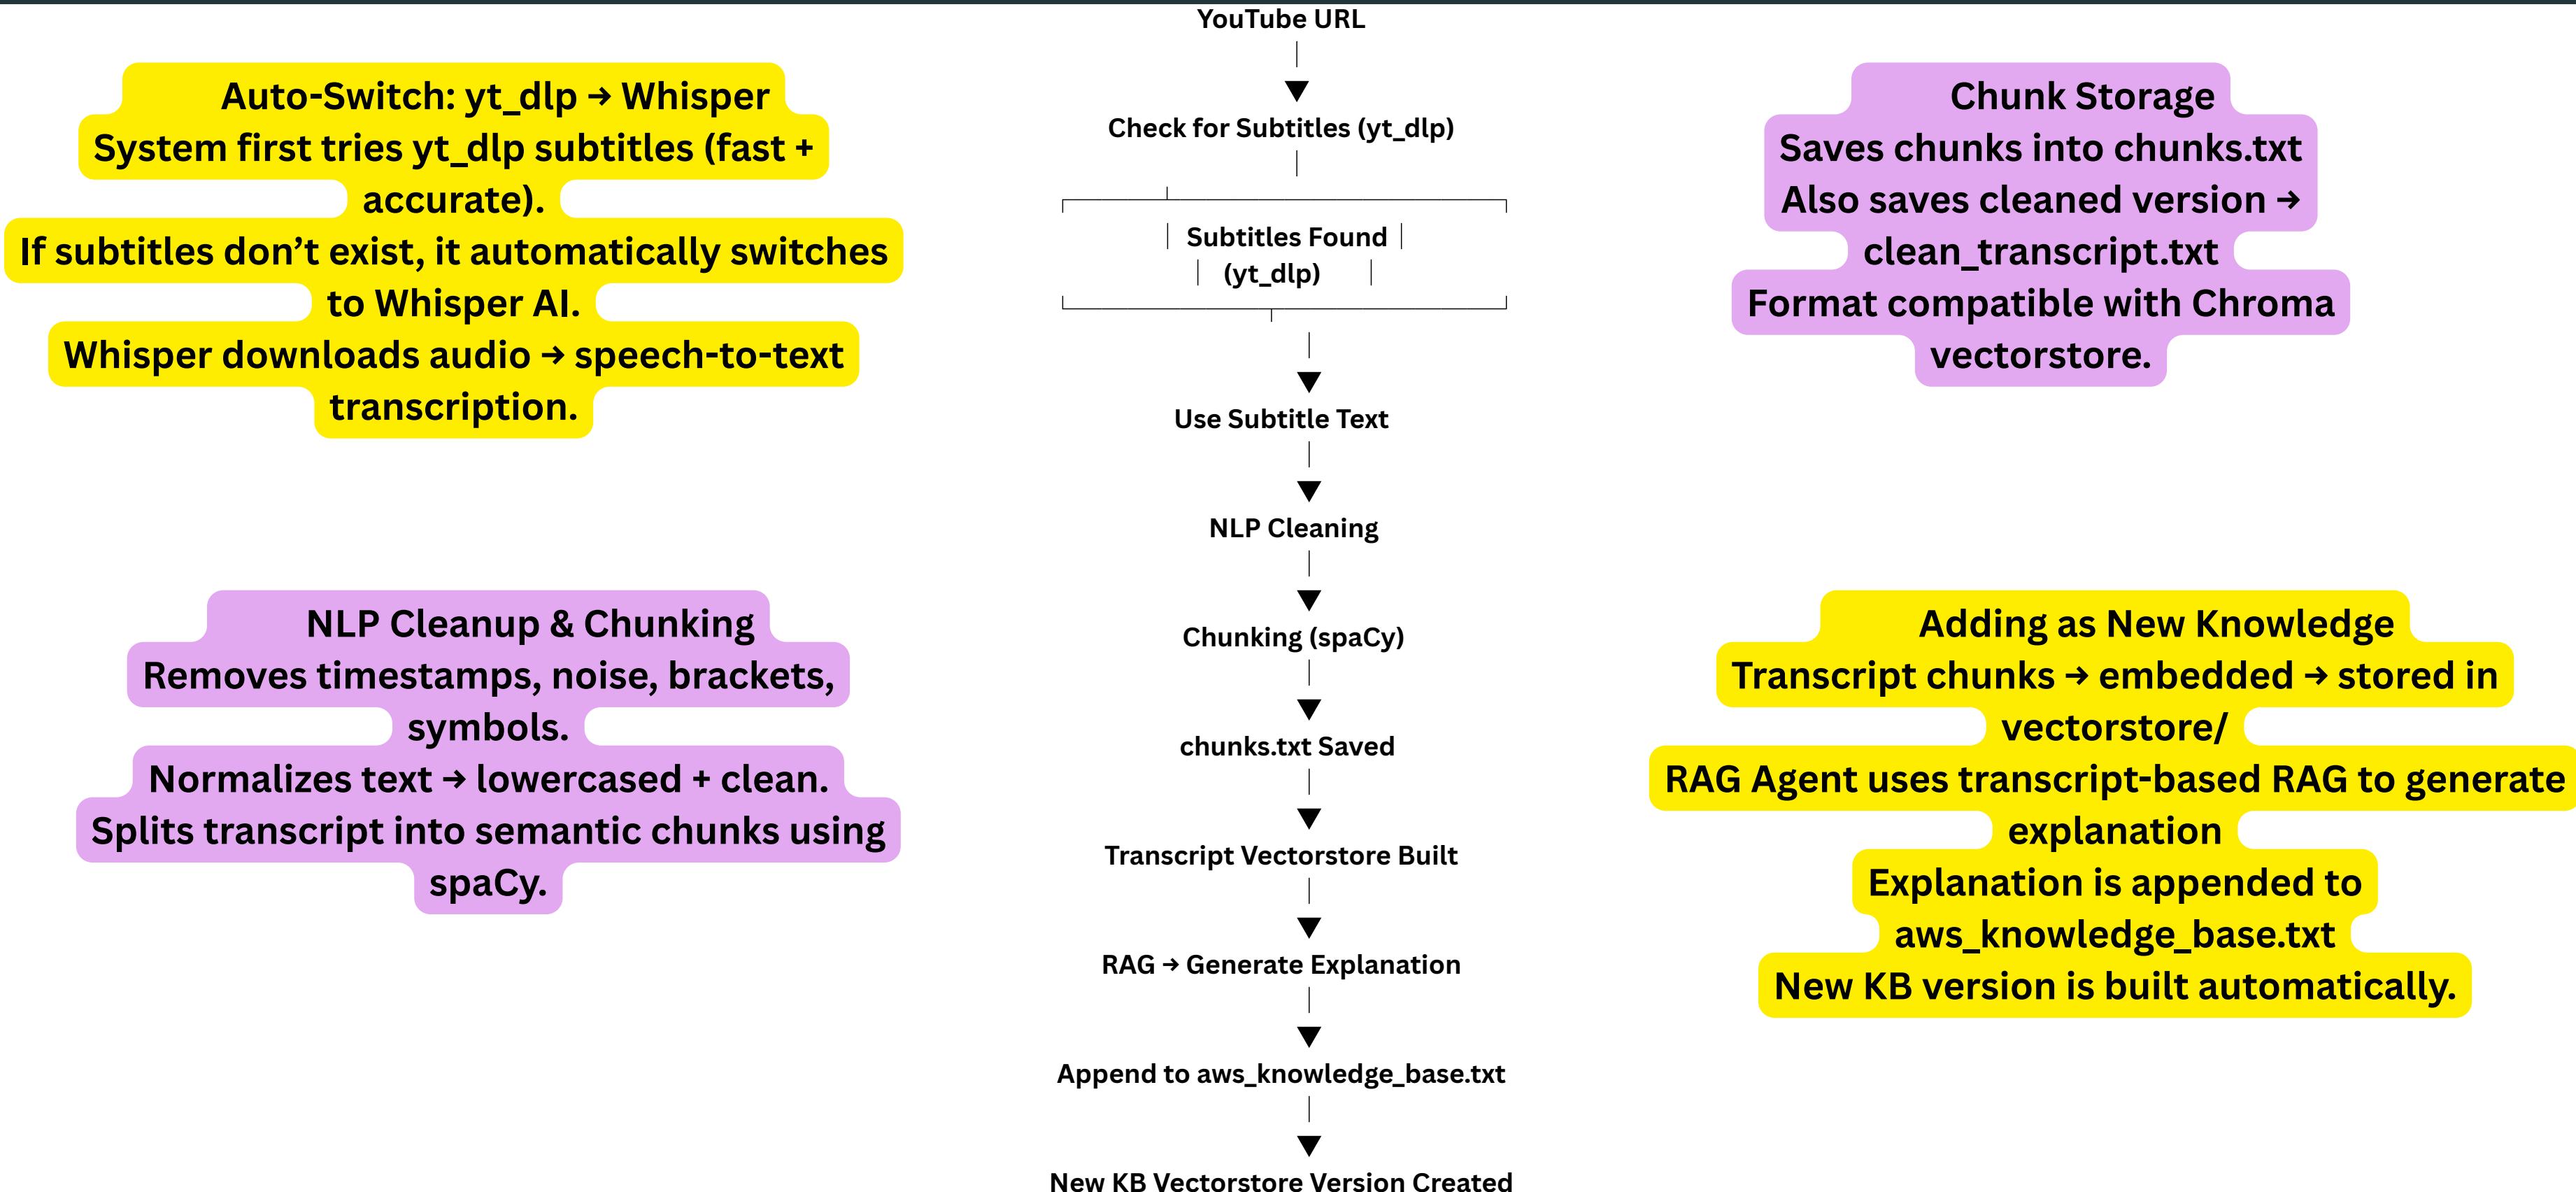
Each mode routes to its own chain:

- teach() → context explanation
- generate_quiz() → MCQs
- generate_code_answer() → code + explanation

Agents summary

Step	What It Does	Type of Agent	Description
Step 4A – Retriever	Finds the most relevant transcript chunks	 Knowledge Agent	Works like a memory retriever – fetches relevant info for the next step.
Step 4B – RAG Chain (Teacher)	Reads retrieved chunks + explains concepts	 Teacher Agent	Uses GPT (ChatOpenAI) to answer beginner questions clearly, step-by-step.
Step 4C-A – Quiz Maker	Creates MCQs & evaluates answers	 Quiz Agent	Generates and grades multiple-choice questions based on transcript context.
Step 4C-B – Code Helper	Generates Python/SQL code + explanations	 Coding Agent	Uses transcript + LLM reasoning to write short code examples with explanation.
Step 4C-C – Role Router	Decides whether to teach, quiz, or code	 Controller / Router Agent	A meta-agent that routes user queries to the correct specialized sub-agent.

YouTube Transcript Extraction & Learning Workflow



How Hallucinations Are Prevented

1. Strict RAG Enforcement

LLM receives context + instruction:

“If the context does NOT contain the answer, say:
I don’t see that in my available knowledge.”

2. AWS KB Relevance Check

If topic not in vectorstore →

LLM is NOT allowed to answer directly.

Instead:

→ enters YouTube Learning Mode

→ learns real content → stores it → avoids hallucinated answers.

3. Controlled Prompt Templates

- TEACH, QUIZ, CODE prompts prevent open-ended generation
- Format instructions (JSON) reduce randomness

4. Separation of Knowledge Sources

- AWS KB vectorstore → authoritative content
- Transcript vectorstore → used only during learning
- LLM outputs depend ONLY on retrieved chunks

5. Topic Normalization

Ensures ambiguous user input (e.g., “glue”) becomes “AWS Glue” → relevant retrieval instead of guessing.